

Lab M1.03 - sklearn Model Training + Evaluation

Anson Knausenberger

Part 1 (Guided Exercise):

Screenshots or output showing:

Dataset exploration

```
=====
BASIC STATISTICS
=====

      mean radius  mean texture  mean perimeter  mean area \
count    569.000000    569.000000    569.000000    569.000000
mean     14.127292    19.289649    91.969033   654.889104
std      3.524049    4.301036    24.298981   351.914129
min      6.981000    9.710000    43.790000   143.500000
25%     11.700000    16.170000    75.170000   420.300000
50%     13.370000    18.840000    86.240000   551.100000
75%     15.780000    21.800000   104.100000  782.700000
max     28.110000    39.280000   188.500000  2501.000000

      mean smoothness  mean compactness  mean concavity  mean concave points \
count    569.000000    569.000000    569.000000    569.000000
mean      0.096360     0.104341     0.088799     0.048919
std       0.014064     0.052813     0.079720     0.038803
min       0.052630     0.019380     0.000000     0.000000
25%      0.086370     0.064920     0.029560     0.020310
50%      0.095870     0.092630     0.061540     0.033500
75%      0.105300     0.130400     0.130700     0.074000
max      0.163400     0.345400     0.426800     0.201200


      mean symmetry  mean fractal dimension ... worst texture \
...
=====

FEATURE DISTRIBUTIONS
=====

Saved visualization to 'feature_distributions.png'

Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

Figure 1



Step 2 Observations

- `df.describe()` shows all 30 features are numeric and on different scales (some features like area/perimeter have much larger ranges than others).
- There are no missing values in the dataset.
- The histograms show clear separation: malignant tumors generally have larger values for mean radius, mean perimeter, and mean area than benign tumors (texture overlaps more).

```
=====
DATA SPLIT
=====
Training set size: 455 samples
Test set size: 114 samples
Training features: 30
Test features: 30

Training set target distribution:
target
1    285
0    170
Name: count, dtype: int64
  Malignant (0): 170 (37.4%)
  Benign (1): 285 (62.6%)

Test set target distribution:
target
1    72
0    42
Name: count, dtype: int64
  Malignant (0): 42 (36.8%)
  Benign (1): 72 (63.2%)
```

Step 3 Observations

- I separated the dataset into features `X` (30 columns) and target `y` (569 labels).
- I split the data into 455 training samples (80%) and 114 test samples (20%) using `random_state=42` for reproducibility.
- Using `stratify=y` kept the malignant/benign proportions nearly the same in both train and test sets.

```
(Bootcamp-env) (Bootcamp-env) PS C:\Users\anson\OneDrive\Desktop\!IronHack\AI Labs\Week 1\Lab M1.03> & C:\Users\anson\anaconda3\envs\Bootcamp-env\python.exe "c:/Users/anson/OneDrive/Desktop/!IronHack/AI Labs/Week 1/Lab M1.03/breast_cancer_prediction.py"
Loading breast cancer dataset...
```

Dataset type: <class 'sklearn.utils._bunch.Bunch'>
Number of samples: 569
Number of features: 30
Target classes: ['malignant' 'benign']

First few rows:

	mean radius	mean texture	mean perimeter	mean area	...	worst concave points	worst symmetry	worst fractal dimension	target
0	17.99	10.38	122.80	1001.0	...	0.2654	0.4601	0.11890	0
1	20.57	17.77	132.90	1326.0	...	0.1860	0.2750	0.08902	0
2	19.69	21.25	130.00	1203.0	...	0.2430	0.3613	0.08758	0
3	11.42	20.38	77.58	386.1	...	0.2575	0.6638	0.17300	0
4	20.29	14.34	135.10	1297.0	...	0.1625	0.2364	0.07678	0

[5 rows x 31 columns]

Dataset info:
<class 'pandas.DataFrame'>
RangeIndex: 569 entries, 0 to 568

Data columns (total 31 columns):

#	Column	Non-Null Count	Dtype
0	mean radius	569	non-null float64
1	mean texture	569	non-null float64
2	mean perimeter	569	non-null float64
3	mean area	569	non-null float64
4	mean smoothness	569	non-null float64
5	mean compactness	569	non-null float64
6	mean concavity	569	non-null float64
7	mean concave points	569	non-null float64
8	mean symmetry	569	non-null float64
9	mean fractal dimension	569	non-null float64
10	radius error	569	non-null float64
11	texture error	569	non-null float64
12	perimeter error	569	non-null float64
13	area error	569	non-null float64
14	smoothness error	569	non-null float64
15	compactness error	569	non-null float64
16	concavity error	569	non-null float64
17	concave points error	569	non-null float64
18	symmetry error	569	non-null float64
19	fractal dimension error	569	non-null float64
20	worst radius	569	non-null float64
21	worst texture	569	non-null float64
22	worst perimeter	569	non-null float64
23	worst area	569	non-null float64
24	worst smoothness	569	non-null float64
25	worst compactness	569	non-null float64
26	worst concavity	569	non-null float64
27	worst concave points	569	non-null float64
28	worst symmetry	569	non-null float64
29	worst fractal dimension	569	non-null float64
30	target	569	non-null int64

dtypes: float64(30), int64(1)

memory usage: 137.9 KB

None

Target distribution:

target

1 357

0 212

Name: count, dtype: int64

Malignant (0): 212

Benign (1): 357

=====

BASIC STATISTICS

```
=====  
 mean radius mean texture mean perimeter ... worst symmetry worst fractal dimension  
target  
count 569.000000 569.000000 569.000000 ... 569.000000 569.000000 569.000000  
mean 14.127292 19.289649 91.969033 ... 0.290076 0.083946 0.627417  
std 3.524049 4.301036 24.298981 ... 0.061867 0.018061 0.483918  
min 6.981000 9.710000 43.790000 ... 0.156500 0.055040 0.000000  
25% 11.700000 16.170000 75.170000 ... 0.250400 0.071460 0.000000  
50% 13.370000 18.840000 86.240000 ... 0.282200 0.080040 1.000000  
75% 15.780000 21.800000 104.100000 ... 0.317900 0.092080 1.000000  
max 28.110000 39.280000 188.500000 ... 0.663800 0.207500 1.000000
```

[8 rows x 31 columns]

```
=====  
MISSING VALUES  
=====
```

✓ No missing values found!

```
=====  
FEATURE DISTRIBUTIONS  
=====
```

Saved visualization to 'feature_distributions.png'

Model training

```
print("KNN classifier trained successfully!")
print(f"Number of neighbors (k): {knn.n_neighbors}")

# Make predictions
y_train_pred = knn.predict(X_train)
y_test_pred = knn.predict(X_test)

print(f"\nTraining predictions: {len(y_train_pred)}")
print(f"Test predictions: {len(y_test_pred)}")
```

Python

```
KNN classifier trained successfully!
Number of neighbors (k): 5
```

```
Training predictions: 455
Test predictions: 114
```

Step 4 Observations

- I trained a K-Nearest Neighbors (KNN) classifier with $k = 5$ using the training dataset.
- KNN predicts a new sample by finding the 5 most similar (closest) training samples and choosing the most common class among them.
- The model produced predictions for all 455 training samples and 114 test samples without errors.

```

=====
STEP 5: PREDICTIONS vs ACTUAL
=====

First 15 predictions vs actual:
   Actual  Predicted Actual_label Predicted_label
0       0         0 malignant    malignant
1       1         1 benign      benign
2       0         0 malignant    malignant
3       1         0 benign      malignant
4       0         0 malignant    malignant
5       1         1 benign      benign
6       1         1 benign      benign
7       0         0 malignant    malignant
8       0         0 malignant    malignant
9       0         0 malignant    malignant
10      1         1 benign      benign
11      0         0 malignant    malignant
12      1         1 benign      benign
13      0         0 malignant    malignant
14      0         0 malignant    malignant

Correct in first 15 examples: 14/15
...
76      0         1 malignant    benign
80      1         0 benign      malignant
99      1         0 benign      malignant
102     0         1 malignant    benign

Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

```

Step 5 Observations

- The model successfully generated predictions for every sample in the test set.
- In the first 15 test examples, 14 predictions matched the actual label.
- There were 10 total mismatches in the full test set, which we'll quantify more formally using evaluation metrics next.

Evaluation metrics

```

==== Model Performance ===

Training Accuracy: 0.9473 (94.73%)
Test Accuracy: 0.9123 (91.23%)

(Considering MALIGNANT as the positive class: target=0)
Test Precision (malignant): 0.8636
Test Recall (malignant): 0.9048

==== Confusion Matrix (rows=Actual, cols=Predicted) ===
Order: 0=malignant, 1=benign
[[38  4]
 [ 6 66]]

Interpretation:
Correct malignant (actual 0 → predicted 0): 38
Missed malignant (actual 0 → predicted 1): 4
Benign flagged malignant (actual 1 → predicted 0): 6
Correct benign (actual 1 → predicted 1): 66

==== Classification Report ===
      precision    recall   f1-score  support
malignant       0.86     0.90     0.88      42
benign         0.94     0.92     0.93      72
...
accuracy           -         -     0.91     114
macro avg        0.90     0.91     0.91     114
weighted avg     0.91     0.91     0.91     114

```

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output settings...

Step 6 Observations

- The model achieved 94.73% training accuracy and 91.23% test accuracy, which suggests decent generalization with some errors on unseen data.
- Treating malignant as the “positive” class (target = 0), precision is 0.86 and recall is 0.90, meaning it correctly identifies most malignant cases but still misses some.
- The confusion matrix shows 4 missed malignant cases (malignant predicted as benign) and 6 false alarms (benign predicted as malignant); missed malignant cases are the most concerning in a medical screening context.

K value comparison results

```
% Generate + Code + Markdown | Run All Clear All Outputs | Outline ...
Select Ke
D v
# Make predictions
y_pred_temp = knn_temp.predict(X_test)

# Calculate metrics
acc = accuracy_score(y_test, y_pred_temp)

# malignant is class 0 in this dataset, so treat it as "positive"
prec_malig = precision_score(y_test, y_pred_temp, pos_label=0)
rec_malig = recall_score(y_test, y_pred_temp, pos_label=0)

results.append({
    'K': k,
    'Accuracy': acc,
    'Precision_malignant': prec_malig,
    'Recall_malignant': rec_malig
})

print(f"K={k:2d}: Accuracy={acc:.4f}, Precision(malig)={prec_malig:.4f}, Recall(malig)={rec_malig:.4f}")

# Find best K by accuracy (simple approach)
results_df = pd.DataFrame(results)
best_k = results_df.loc[results_df['Accuracy'].idxmax(), 'K']
print(f"\nBest K value by Accuracy: {best_k} (Accuracy: {results_df['Accuracy'].max():.4f})")

# Visualize results
plt.figure(figsize=(10, 6))
plt.plot(results_df['K'], results_df['Accuracy'], marker='o', label='Accuracy')
plt.plot(results_df['K'], results_df['Precision_malignant'], marker='s', label='Precision (malignant)')
plt.plot(results_df['K'], results_df['Recall_malignant'], marker='^', label='Recall (malignant)')
plt.xlabel('K (Number of Neighbors)')
plt.ylabel('Score')
plt.title('KNN Performance vs K Value')
plt.legend()
plt.grid(True, alpha=0.3)
plt.savefig('knn_k_comparison.png', dpi=150, bbox_inches='tight')
print("\nSaved visualization to 'knn_k_comparison.png'")
plt.show()


results_df
```

```
=====
EXPERIMENTING WITH DIFFERENT K VALUES
=====

K= 1: Accuracy=0.9211, Precision(malig)=0.8667, Recall(malig)=0.9286
K= 3: Accuracy=0.9298, Precision(malig)=0.9048, Recall(malig)=0.9048
K= 5: Accuracy=0.9123, Precision(malig)=0.8636, Recall(malig)=0.9048
K= 7: Accuracy=0.9298, Precision(malig)=0.9048, Recall(malig)=0.9048
K= 9: Accuracy=0.9386, Precision(malig)=0.9268, Recall(malig)=0.9048
K=11: Accuracy=0.9386, Precision(malig)=0.9268, Recall(malig)=0.9048
```

Best K value by Accuracy: 9 (Accuracy: 0.9386)

Saved visualization to 'knn_k_comparison.png'



K	Accuracy	Precision_malignant	Recall_malignant
0	0.921053	0.866667	0.928571
1	0.929825	0.904762	0.904762
2	0.912281	0.863636	0.904762
3	0.929825	0.904762	0.904762
4	0.938596	0.926829	0.904762
5	0.938596	0.926829	0.904762


Step 7 Observations

- Changing K changes performance: smaller K can be more sensitive (higher recall) but may create more false alarms, while larger K smooths decisions and can improve stability.
- In my results, K=9 (and K=11) achieved the highest test accuracy (~0.939) and the highest precision for malignant (~0.927), while recall for malignant stayed ~0.905.
- If the priority is to catch as many malignant cases as possible, K=1 had the highest malignant recall (~0.929), but with lower precision (more false positives).

Part 2 (Unguided Exercise):

Screenshots or output showing:

- Data exploration and preprocessing
- Model training and evaluation
- Comparison of different K values



```
=====
STEP 2: PREPROCESSING
=====
Missing values after TotalCharges conversion:
TotalCharges      11
gender            0
Partner           0
SeniorCitizen     0
Dependents        0
tenure            0
MultipleLines     0
PhoneService      0
OnlineSecurity    0
OnlineBackup       0
dtype: int64

Encoded shape (rows, columns): (7043, 31)
X shape: (7043, 30)
y shape: (7043,)

Target distribution after encoding:
Churn
0    5174
1    1869
Name: count, dtype: int64
Churn rate (%): 26.54
```

```
=====
STEP 3: TRAIN/TEST SPLIT
=====
Train size: 5634
Test size: 1409

Train churn distribution:
Churn
0    4139
1    1495
Name: count, dtype: int64
Train churn rate (%): 26.54

Test churn distribution:
Churn
0    1035
1    374
Name: count, dtype: int64
Test churn rate (%): 26.54
```

```
=====  
STEP 4: TRAIN KNN (k=5)  
=====
```

```
Trained KNN with k=5  
Train predictions: 5634  
Test predictions: 1409
```

```
=====  
STEP 5: EVALUATION (positive = churn=1)  
=====
```

```
Training Accuracy: 0.8277 (82.77%)  
Test Accuracy: 0.7658 (76.58%)  
Test Precision (churn): 0.5791  
Test Recall (churn): 0.4305
```

```
Confusion Matrix (rows=Actual, cols=Predicted)
```

```
Order: 0=No churn, 1=Churn
```

```
[[918 117]  
 [213 161]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
No	0.81	0.89	0.85	1035
Yes	0.58	0.43	0.49	374
accuracy			0.77	1409
macro avg	0.70	0.66	0.67	1409
weighted avg	0.75	0.77	0.75	1409

```
(Bootcamp-env) (Bootcamp-env) PS C:\Users\anson\OneDrive\Desktop\!IronHack\AI Labs\Week 1\Lab M1.03>
```

```

=====
STEP 6: K COMPARISON
=====
K= 1: Accuracy=0.7119, Precision=0.4570, Recall=0.4545
K= 3: Accuracy=0.7622, Precision=0.5657, Recall=0.4492
K= 5: Accuracy=0.7658, Precision=0.5791, Recall=0.4305
K= 7: Accuracy=0.7814, Precision=0.6260, Recall=0.4385
K= 9: Accuracy=0.7885, Precision=0.6532, Recall=0.4332
K=11: Accuracy=0.7871, Precision=0.6623, Recall=0.4037
K=15: Accuracy=0.7885, Precision=0.6681, Recall=0.4037

Results table:
   K  Accuracy  Precision  Recall
0  1  0.711852  0.456989  0.454545
1  3  0.762243  0.565657  0.449198
2  5  0.765791  0.579137  0.430481
3  7  0.781405  0.625954  0.438503
4  9  0.788502  0.653226  0.433155
5 11  0.787083  0.662281  0.403743
6 15  0.788502  0.668142  0.403743

Best K by Recall (catching churners): 1
Best K by Accuracy: 9

```

A brief report (1-2 pages) including:

Summary of approach

- Model performance metrics
- Key findings about customer churn
- Business recommendations
- Limitations and future improvements

Telco Customer Churn Prediction (KNN) — Brief Report

1) Summary of approach

I built a supervised classification model to predict whether a telecom customer will churn (Yes/No) using the Telco Customer Churn dataset (7,043 customers; 21 columns). I followed a standard ML workflow: load and explore the dataset, preprocess/encode features, split into training and test sets (80/20 with stratification), train a K-Nearest Neighbors (KNN) classifier, evaluate using accuracy/precision/recall + confusion matrix, and then tune the number of neighbors (K). Because KNN is distance-based and sensitive to feature scale, I also implemented an improved sklearn Pipeline that standardizes numeric features and one-hot encodes categoricals in a single reproducible workflow.

2) Data exploration (what I observed)

- The target variable is imbalanced: **Churn = Yes: 1,869 (26.54%), Churn = No: 5,174 (73.46%).**

- TotalCharges was loaded as a string; converting it to numeric revealed **11 missing/blank values**, which were filled using the median.
- A quick histogram comparison suggested **churners tend to have higher MonthlyCharges** (the churn “Yes” distribution is more concentrated at higher monthly costs), even though the total count of churners is lower.

3) Preprocessing

- Dropped customerID (identifier, not predictive).
- Converted TotalCharges to numeric and filled missing values with the median.
- Converted the target Churn to binary: **No = 0, Yes = 1**.
- Encoded categorical variables using one-hot encoding (either manually via get_dummies or via OneHotEncoder in a Pipeline).
- Train/test split: **Train = 5,634, Test = 1,409**, with identical churn rate (**26.54%**) due to stratification.

4) Model performance metrics

Baseline KNN (manual one-hot, no scaling)

Using KNN with K=5:

- **Test Accuracy: 0.7658**
- **Precision (Churn=Yes): 0.5791**
- **Recall (Churn=Yes): 0.4305**
- **Confusion matrix** (rows=actual, cols=pred; 0=No, 1=Yes):

```
[  
\begin{bmatrix}  
918 & 117 \\\  
213 & 161  
\end{bmatrix}]
```

Interpretation: the model correctly flagged **161 churners**, but **missed 213 churners** (low recall).

K sweep (no scaling) showed:

- Best recall was **K=1** with **Recall ~0.4545**
- Best accuracy was **K=9/15** with **Accuracy ~0.7885** (but recall stayed ~0.40–0.43)

Improved KNN (Pipeline: scaling numeric + one-hot encoding categoricals)

With a proper preprocessing Pipeline, performance improved substantially. Best results were at

K=9:

- **Accuracy: 0.7800**
- **Precision (Churn=Yes): 0.5856**
- **Recall (Churn=Yes): 0.5856**

This is a major improvement in recall (catching churners) versus the unscaled version (~0.43–0.46 recall).

5) Key findings about customer churn

- **Churn is meaningfully imbalanced** (~26.5% churn), so accuracy alone is not sufficient; recall/precision for churners is more important for retention actions.
- **MonthlyCharges appears associated with churn** (churners skew toward higher monthly costs). This is consistent with the idea that higher monthly bills can increase churn risk, especially if perceived value is low.
- **Preprocessing quality strongly affects KNN**, especially scaling numeric features. The pipeline approach improved churn recall from ~0.43 to ~0.59.

6) Business recommendations

1. **Use the pipeline KNN (K=9) as the best baseline** from this exercise. It catches ~59% of churners, which is meaningfully better than the unscaled KNN.
2. **Operationalize around recall vs precision tradeoffs:**
 - o If the business wants to catch more churners (maximize retention saves), prioritize **recall** even if precision drops (more outreach).
 - o If outreach capacity is limited, optimize for a better precision/recall balance and target the highest-risk segment only.
3. **Target interventions where the model suggests higher churn risk**, especially customers with **higher MonthlyCharges**. Possible actions: proactive discounts, service upgrades, contract incentives, improved support, or “value reinforcement” messaging.
4. **Next analytics step (quick wins):** segment churn rates by Contract type, tenure bands, PaymentMethod, and InternetService to identify the strongest levers for retention campaigns.

7) Limitations and future improvements

- **Model limitation (KNN):** KNN can struggle with high-dimensional one-hot encoded data and is sensitive to scaling; while the pipeline helps, KNN may still be suboptimal for this dataset.
- **Metric limitation:** We did not optimize thresholds (KNN outputs class labels by default). For retention, probability scores + threshold tuning are often better.
- **Feature importance:** KNN is not easily interpretable for feature importance. For clearer drivers of churn, consider more interpretable models.
- **Future improvements:**
 1. Try **Logistic Regression** (interpretable coefficients), **Random Forest**, or **Gradient Boosting** and compare recall/precision/F1 for churners.
 2. Use **cross-validation** and a more systematic hyperparameter search.
 3. Evaluate **ROC-AUC / PR-AUC** and tune decision thresholds based on business cost of false negatives vs false positives.
 4. Consider handling class imbalance explicitly (e.g., class weights in other models, resampling methods).